

```
In [1]: import numpy as np
import pandas as pd
```

```
In [3]: from sklearn.datasets import load_boston
boston = load_boston()
```

```
In [4]: data = pd.DataFrame(boston.data)
```

```
In [5]: data.head()
```

```
Out[5]:
```

	0	1	2	3	4	5	6	7	8	9	10	11	12
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33

```
In [7]: data.columns = boston.feature_names
```

```
In [8]: data['PRICE'] = boston.target
```

```
In [9]: data.head(n=10)
```

```
Out[9]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LST
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33
5	0.02985	0.0	2.18	0.0	0.458	6.430	58.7	6.0622	3.0	222.0	18.7	394.12	5.02
6	0.08829	12.5	7.87	0.0	0.524	6.012	66.6	5.5605	5.0	311.0	15.2	395.60	12.63
7	0.14455	12.5	7.87	0.0	0.524	6.172	96.1	5.9505	5.0	311.0	15.2	396.90	19.63
8	0.21124	12.5	7.87	0.0	0.524	5.631	100.0	6.0821	5.0	311.0	15.2	386.63	29.76
9	0.17004	12.5	7.87	0.0	0.524	6.004	85.9	6.5921	5.0	311.0	15.2	386.71	17.02

```
In [10]: print(data.shape)
```

```
(506, 14)
```

```
In [11]: data.isnull().sum()
```


```
Out[11]: CRIM      0
          ZN        0
          INDUS    0
          CHAS     0
          NOX      0
          RM       0
          AGE      0
          DIS      0
          RAD      0
          TAX      0
          PTRATIO  0
          B        0
          LSTAT    0
          PRICE    0
          dtype: int64
```

```
In [12]: data.describe
```

```
Out[12]: <bound method NDFrame.describe of
RM      AGE      DIS      RAD      TAX      \
0      0.00632  18.0    2.31    0.0    0.538  6.575  65.2  4.0900  1.0  296.0
1      0.02731   0.0    7.07    0.0    0.469  6.421  78.9  4.9671  2.0  242.0
2      0.02729   0.0    7.07    0.0    0.469  7.185  61.1  4.9671  2.0  242.0
3      0.03237   0.0    2.18    0.0    0.458  6.998  45.8  6.0622  3.0  222.0
4      0.06905   0.0    2.18    0.0    0.458  7.147  54.2  6.0622  3.0  222.0
5      0.02985   0.0    2.18    0.0    0.458  6.430  58.7  6.0622  3.0  222.0
6      0.08829  12.5    7.87    0.0    0.524  6.012  66.6  5.5605  5.0  311.0
7      0.14455  12.5    7.87    0.0    0.524  6.172  96.1  5.9505  5.0  311.0
8      0.21124  12.5    7.87    0.0    0.524  5.631 100.0  6.0821  5.0  311.0
9      0.17004  12.5    7.87    0.0    0.524  6.004  85.9  6.5921  5.0  311.0
10     0.22489  12.5    7.87    0.0    0.524  6.377  94.3  6.3467  5.0  311.0
11     0.11747  12.5    7.87    0.0    0.524  6.009  82.9  6.2267  5.0  311.0
12     0.09378  12.5    7.87    0.0    0.524  5.889  39.0  5.4509  5.0  311.0
13     0.62976   0.0    8.14    0.0    0.538  5.949  61.8  4.7075  4.0  307.0
14     0.63796   0.0    8.14    0.0    0.538  6.096  84.5  4.4619  4.0  307.0
15     0.62739   0.0    8.14    0.0    0.538  5.834  56.5  4.4986  4.0  307.0
16     1.05393   0.0    8.14    0.0    0.538  5.935  29.3  4.4986  4.0  307.0
17     0.78420   0.0    8.14    0.0    0.538  5.000  81.7  4.3570  4.0  307.0
```


```
In [16]: import seaborn as sns  
sns.distplot(data.PRICE)
```

Out[16]: <matplotlib.axes._subplots.AxesSubplot at 0x7f1096facef0>



```
In [17]: sns.boxplot(data.PRICE)
```

Out[17]: <matplotlib.axes._subplots.AxesSubplot at 0x7f1094edf4a8>



```
In [18]: correlation = data.corr()  
correlation.loc['PRICE']
```

```
Out[18]: CRIM      -0.388305  
          ZN        0.360445  
          INDUS    -0.483725  
          CHAS      0.175260  
          NOX      -0.427321  
          RM        0.695360  
          AGE      -0.376955  
          DIS       0.249929  
          RAD      -0.381626  
          TAX      -0.468536  
          PTRATIO  -0.507787  
          B         0.333461  
          LSTAT    -0.737663  
          PRICE     1.000000  
          Name: PRICE, dtype: float64
```

```
In [19]: import matplotlib.pyplot as plt
```

```
In [20]: fig, axes = plt.subplots(figsize=(15,12))
sns.heatmap(correlation, square = True, annot = True)
```

```
Out[20]: <matplotlib.axes._subplots.AxesSubplot at 0x7f1094e57438>
```



```
In [24]: plt.figure(figsize = (20,5))
features = ['LSTAT' , 'RM', 'PTRATIO']
for i, col in enumerate(features):
    plt.subplot(1, len(features), i+1)
    x = data[col]
    y = data.PRICE
    plt.scatter(x,y, marker='o')
    plt.title("Variation in House prices")
    plt.xlabel(col)
    plt.ylabel("house prices in $1000")
```



```
In [27]: x = data.iloc[:, :-1]
y = data.PRICE
```

```
In [35]: from sklearn.model_selection import train_test_split
xtrain, xtest, ytrain, ytest = train_test_split(x, y, test_size=0.3, random_st
```

```
In [37]: mean = xtrain.mean(axis=0)
std = xtrain.std(axis=0)
```

```
In [39]: from sklearn.linear_model import LinearRegression
```

```
In [40]: regressor = LinearRegression()
```

```
In [41]: regressor.fit(xtrain, ytrain)
```

```
Out[41]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,
normalize=False)
```

```
In [42]: ypred = regressor.predict(xtest)
```

```
In [43]: from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
```

```
In [45]: rmse = (np.sqrt(mean_squared_error(ytest, ypred)))
rmse
```

```
Out[45]: 5.214975145375418
```

```
In [47]: r2 = r2_score(ytest, ypred)
r2
```

```
Out[47]: 0.6733825506400176
```

```
In [88]: from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
xtrain = sc.fit_transform(xtrain)
xtest = sc.transform(xtest)
```

```
In [89]: import keras
```

```
In [90]: from keras.layers import Dense, Activation, Dropout
from keras.models import Sequential
```

```
In [91]: model = Sequential()
model.add(Dense(128, activation='relu', input_dim=13))
model.add(Dense(64, activation='relu'))
model.add(Dense(32, activation='relu'))
model.add(Dense(16, activation='relu'))
model.add(Dense(1))
```

```
In [93]: model.compile(optimizer = 'adam', loss = 'mean_squared_error', metrics=['mae'])
```

```
In [94]: !pip install ann_visualizer
```

Requirement already satisfied: ann_visualizer in /home/rmdstic/anaconda3/lib/python3.7/site-packages (2.5)

```
In [95]: !pip install graphviz
```

Requirement already satisfied: graphviz in /home/rmdstic/anaconda3/lib/python3.7/site-packages (0.20.1)

```
In [96]: from ann_visualizer.visualize import ann_viz;
```

```
In [97]: ann_viz(model, title="DEMO ANN");
```



```
In [105]: fig = go.Figure()
fig.add_trace(go.Scattergl(y=history.history['loss'], name='Train'))
fig.add_trace(go.Scattergl(y=history.history['val_loss'], name='Vaild'))
fig.update_layout(height=500, width=700, xaxis_title = 'Epoch', yaxis_title='L
fig.show()
```



```
In [106]: fig = go.Figure()
fig.add_trace(go.Scattergl(y=history.history['mae'], name='Train'))
fig.add_trace(go.Scattergl(y=history.history['val_mae'], name='Valid'))
fig.update_layout(height=500, width=700, xaxis_title = 'Epoch', yaxis_title='M
fig.show()
```



```
In [107]: y_pred = model.predict(xtest)
```

```
In [108]: mse_nn, mae_nn = model.evaluate(xtest, ytest)
print('Mean Squared Error on test data: ', mse_nn)
print('Mean Squared Error on test data: ', mae_nn)
```

```
5/5 [=====] - 0s 1ms/step - loss: 7002.5273 - mae: 7
1.1850
Mean Squared Error on test data: 7002.52734375
Mean Squared Error on test data: 71.18496704101562
```

```
In [109]: from sklearn.metrics import mean_absolute_error
```

```
In [110]: lr_model = LinearRegression()
```

```
In [111]: lr_model.fit(xtrain, ytrain)
```

```
Out[111]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None,  
normalize=False)
```

```
In [112]: y_pred_lr = lr_model.predict(xtest)
```

```
In [113]: mse_lr = mean_squared_error(ytest, y_pred_lr)  
mae_lr = mean_absolute_error(ytest, y_pred_lr)
```

```
In [114]: print('Mean Squared Error on Test Data: ', mse_lr)  
print('Mean Absolute Error on Test Data: ', mae_lr)
```

```
Mean Squared Error on Test Data: 454.4737836456808  
Mean Absolute Error on Test Data: 18.963407281929687
```

```
In [115]: from sklearn.metrics import r2_score
```

```
In [121]: r2 = r2_score(ytest, ypred)  
print(r2)
```

```
0.6733825506400176
```

```
In [122]: from sklearn.metrics import mean_squared_error
```

```
In [124]: rmse = np.sqrt(mean_squared_error(ytest, ypred))  
print(rmse)
```

```
5.214975145375418
```

```
In [128]: import sklearn  
new_data = sklearn.preprocessing.StandardScaler().fit_transform([[0.1,10.0, 5
```

```
In [129]: prediction = model.predict(new_data)
```

```
In [131]: print('Prdeicted house price: ', prediction)
```

```
Prdeicted house price: [[9.389934]]
```

```
In [ ]:
```

```
In [ ]:
```

