

```
In [1]: # Akshata NLP_04
```

```
In [2]: # Create a transformer from scratch using the Pytorch Library.
```

```
In [8]: pip install torch
```

```
Requirement already satisfied: torch in c:\users\tech bazaar\anaconda3\lib\site-packages (2.1.2)
Requirement already satisfied: typing-extensions in c:\users\tech bazaar\anaconda3\lib\site-packages (from torch) (4.9.0)
Requirement already satisfied: filelock in c:\users\tech bazaar\anaconda3\lib\site-packages (from torch) (3.3.1)
Requirement already satisfied: fsspec in c:\users\tech bazaar\anaconda3\lib\site-packages (from torch) (2021.10.1)
Requirement already satisfied: networkx in c:\users\tech bazaar\anaconda3\lib\site-packages (from torch) (2.6.3)
Requirement already satisfied: Jinja2 in c:\users\tech bazaar\anaconda3\lib\site-packages (from torch) (2.11.3)
Requirement already satisfied: sympy in c:\users\tech bazaar\anaconda3\lib\site-packages (from torch) (1.9)
Requirement already satisfied: MarkupSafe>=0.23 in c:\users\tech bazaar\anaconda3\lib\site-packages (from Jinja2->torch) (1.1.1)
Requirement already satisfied: mpmath>=0.19 in c:\users\tech bazaar\anaconda3\lib\site-packages (from sympy->torch) (1.2.1)
Note: you may need to restart the kernel to use updated packages.
```

```
In [9]: pip install torchtext
```

```
Collecting torchtext
  Downloading torchtext-0.16.2-cp39-cp39-win_amd64.whl (1.9 MB)
Requirement already satisfied: torch==2.1.2 in c:\users\tech bazaar\anaconda3\lib\site-packages (from torchtext) (2.1.2)
Requirement already satisfied: tqdm in c:\users\tech bazaar\anaconda3\lib\site-packages (from torchtext) (4.62.3)
Requirement already satisfied: numpy in c:\users\tech bazaar\anaconda3\lib\site-packages (from torchtext) (1.22.4)
Requirement already satisfied: requests in c:\users\tech bazaar\anaconda3\lib\site-packages (from torchtext) (2.26.0)
Collecting torchdata==0.7.1
  Downloading torchdata-0.7.1-cp39-cp39-win_amd64.whl (1.3 MB)
Requirement already satisfied: networkx in c:\users\tech bazaar\anaconda3\lib\site-packages (from torch==2.1.2->torchtext) (2.6.3)
Requirement already satisfied: sympy in c:\users\tech bazaar\anaconda3\lib\site-packages (from torch==2.1.2->torchtext) (1.9)
Requirement already satisfied: Jinja2 in c:\users\tech bazaar\anaconda3\lib\site-packages (from torch==2.1.2->torchtext) (2.11.3)
Requirement already satisfied: typing-extensions in c:\users\tech bazaar\anaconda3\lib\site-packages (from torch==2.1.2->torchtext) (4.9.0)
Requirement already satisfied: fsspec in c:\users\tech bazaar\anaconda3\lib\site-packages (from torch==2.1.2->torchtext) (2021.10.1)
Requirement already satisfied: filelock in c:\users\tech bazaar\anaconda3\lib\site-packages (from torch==2.1.2->torchtext) (3.3.1)
Requirement already satisfied: urllib3>=1.25 in c:\users\tech bazaar\anaconda3\lib\site-packages (from torchdata==0.7.1->torchtext) (1.26.7)
Requirement already satisfied: MarkupSafe>=0.23 in c:\users\tech bazaar\anaconda3\lib\site-packages (from Jinja2->torch==2.1.2->torchtext) (1.1.1)
Requirement already satisfied: idna<4,>=2.5 in c:\users\tech bazaar\anaconda3\lib\site-packages (from requests->torchtext) (3.2)
Requirement already satisfied: charset-normalizer~=2.0.0 in c:\users\tech bazaar\anaconda3\lib\site-packages (from requests->torchtext) (2.0.4)
Requirement already satisfied: certifi>=2017.4.17 in c:\users\tech bazaar\anaconda3\lib\site-packages (from requests->torchtext) (2020.6.20)
Requirement already satisfied: mpmath>=0.19 in c:\users\tech bazaar\anaconda3\lib\site-packages (from sympy->torch==2.1.2->torchtext) (1.2.1)
Requirement already satisfied: colorama in c:\users\tech bazaar\anaconda3\lib\site-packages (from tqdm->torchtext) (0.4.6)
Installing collected packages: torchdata, torchtext
Successfully installed torchdata-0.7.1 torchtext-0.16.2
Note: you may need to restart the kernel to use updated packages.
```

```
In [10]: import torch.nn as nn
import torch
import torch.nn.functional as F
import math,copy,re
import warnings
import pandas as pd
import numpy as np
import seaborn as sns
import torchtext
import matplotlib.pyplot as plt
warnings.simplefilter("ignore")
print(torch.__version__)
```

```
2.1.2+cpu
```

```
In [11]: class Embedding(nn.Module):
def __init__(self, vocab_size, embed_dim):
    """
    Args:
        vocab_size: size of vocabulary
        embed_dim: dimension of embeddings
    """
    super(Embedding, self).__init__()
    self.embed = nn.Embedding(vocab_size, embed_dim)

def forward(self, x):
    """
    Args:
        x: input vector
    Returns:
        out: embedding vector
    """
    out = self.embed(x)
    return out
```

```

In [12]: class PositionalEmbedding(nn.Module):
def __init__(self,max_seq_len,embed_model_dim):
    """
    Args:
        seq_len: length of input sequence
        embed_model_dim: demension of embedding
    """
    super(PositionalEmbedding, self).__init__()
    self.embed_dim = embed_model_dim

    pe = torch.zeros(max_seq_len,self.embed_dim)
    for pos in range(max_seq_len):
        for i in range(0,self.embed_dim,2):
            pe[pos, i] = math.sin(pos / (10000 ** ((2 * i)/self.embed_dim)))
            pe[pos, i + 1] = math.cos(pos / (10000 ** ((2 * (i + 1))/self.embed_dim)))
    pe = pe.unsqueeze(0)
    self.register_buffer('pe', pe)

def forward(self, x):
    """
    Args:
        x: input vector
    Returns:
        x: output
    """

    # make embeddings relatively larger
    x = x * math.sqrt(self.embed_dim)
    #add constant to embedding
    seq_len = x.size(1)
    x = x + torch.autograd.Variable(self.pe[:, :seq_len], requires_grad=False)
    return x

```

```

In [13]: class MultiHeadAttention(nn.Module):
def __init__(self, embed_dim=512, n_heads=8):
    """
    Args:
        embed_dim: dimension of embedding vector output
        n_heads: number of self attention heads
    """
    super(MultiHeadAttention, self).__init__()

    self.embed_dim = embed_dim    #512 dim
    self.n_heads = n_heads    #8
    self.single_head_dim = int(self.embed_dim / self.n_heads)    #512/8 = 64 . each key,query, value will be of 64d

    #key,query and value matrixes    #64 x 64
    self.query_matrix = nn.Linear(self.single_head_dim, self.single_head_dim, bias=False) # single key matrix for all 8 keys #512x512
    self.key_matrix = nn.Linear(self.single_head_dim, self.single_head_dim, bias=False)
    self.value_matrix = nn.Linear(self.single_head_dim, self.single_head_dim, bias=False)
    self.out = nn.Linear(self.n_heads*self.single_head_dim, self.embed_dim)

def forward(self, key, query, value, mask=None):    #batch_size x sequence_length x embedding_dim    # 32 x 10 x 512
    """
    Args:
        key : key vector
        query : query vector
        value : value vector
        mask: mask for decoder

    Returns:
        output vector from multihead attention
    """
    batch_size = key.size(0)
    seq_length = key.size(1)

    # query dimension can change in decoder during inference.
    # so we cant take general seq_length
    seq_length_query = query.size(1)

    # 32x10x512
    key = key.view(batch_size, seq_length, self.n_heads, self.single_head_dim) #batch_size x sequence_length x n_heads x single_head_dim = (32x10x8x64)
    query = query.view(batch_size, seq_length_query, self.n_heads, self.single_head_dim) #(32x10x8x64)
    value = value.view(batch_size, seq_length, self.n_heads, self.single_head_dim) #(32x10x8x64)

    k = self.key_matrix(key)    # (32x10x8x64)
    q = self.query_matrix(query)
    v = self.value_matrix(value)

    q = q.transpose(1,2) # (batch_size, n_heads, seq_len, single_head_dim)    # (32 x 8 x 10 x 64)
    k = k.transpose(1,2) # (batch_size, n_heads, seq_len, single_head_dim)
    v = v.transpose(1,2) # (batch_size, n_heads, seq_len, single_head_dim)

    # computes attention
    # adjust key for matrix multiplication
    k_adjussted = k.transpose(-1,-2) #(batch_size, n_heads, single_head_dim, seq_len) #(32 x 8 x 64 x 10)
    product = torch.matmul(q, k_adjussted) #(32 x 8 x 10 x 64) x (32 x 8 x 64 x 10) = #(32x8x10x10)

    # fill those positions of product matrix as (-1e20) where mask positions are 0
    if mask is not None:
        product = product.masked_fill(mask == 0, float("-1e20"))

    #dividing by square root of key dimension
    product = product / math.sqrt(self.single_head_dim) # / sqrt(64)

    #applying softmax
    scores = F.softmax(product, dim=-1)

    #mutiply with value matrix
    scores = torch.matmul(scores, v) ##(32x8x 10x 10) x (32 x 8 x 10 x 64) = (32 x 8 x 10 x 64)

    #concatenated output
    concat = scores.transpose(1,2).contiguous().view(batch_size, seq_length_query, self.single_head_dim*self.n_heads) # (32x8x10x64) -> (32x10x8x64)
    output = self.out(concat) #(32,10,512) -> (32,10,512)

    return output

```

```

In [14]: class TransformerBlock(nn.Module):
def __init__(self, embed_dim, expansion_factor=4, n_heads=8):
    super(TransformerBlock, self).__init__()

    """
    Args:
        embed_dim: dimension of the embedding
        expansion_factor: factor which determines output dimension of linear layer
        n_heads: number of attention heads
    """

    self.attention = MultiHeadAttention(embed_dim, n_heads)

    self.norm1 = nn.LayerNorm(embed_dim)
    self.norm2 = nn.LayerNorm(embed_dim)

    self.feed_forward = nn.Sequential(
        nn.Linear(embed_dim, expansion_factor*embed_dim),
        nn.ReLU(),
        nn.Linear(expansion_factor*embed_dim, embed_dim)
    )

    self.dropout1 = nn.Dropout(0.2)
    self.dropout2 = nn.Dropout(0.2)

def forward(self, key, query, value):
    """
    Args:
        key: key vector
        query: query vector
        value: value vector
        norm2_out: output of transformer block
    """

    attention_out = self.attention(key, query, value) #32x10x512
    attention_residual_out = attention_out + value #32x10x512
    norm1_out = self.dropout1(self.norm1(attention_residual_out)) #32x10x512

    feed_fwd_out = self.feed_forward(norm1_out) #32x10x512 -> #32x10x2048 -> 32x10x512
    feed_fwd_residual_out = feed_fwd_out + norm1_out #32x10x512
    norm2_out = self.dropout2(self.norm2(feed_fwd_residual_out)) #32x10x512

    return norm2_out

class TransformerEncoder(nn.Module):
    """
    Args:
        seq_len : length of input sequence
        embed_dim: dimension of embedding
        num_layers: number of encoder layers
        expansion_factor: factor which determines number of linear layers in feed forward layer
        n_heads: number of heads in multihead attention

    Returns:
        out: output of the encoder
    """
def __init__(self, seq_len, vocab_size, embed_dim, num_layers=2, expansion_factor=4, n_heads=8):
    super(TransformerEncoder, self).__init__()

    self.embedding_layer = Embedding(vocab_size, embed_dim)
    self.positional_encoder = PositionalEncoding(seq_len, embed_dim)

    self.layers = nn.ModuleList([TransformerBlock(embed_dim, expansion_factor, n_heads) for i in range(num_layers)])

def forward(self, x):
    embed_out = self.embedding_layer(x)
    out = self.positional_encoder(embed_out)
    for layer in self.layers:
        out = layer(out, out, out)

    return out #32x10x512

```

```

In [15]: class DecoderBlock(nn.Module):
def __init__(self, embed_dim, expansion_factor=4, n_heads=8):
    super(DecoderBlock, self).__init__()

    """
    Args:
        embed_dim: dimension of the embedding
        expansion_factor: factor which determines output dimension of linear layer
        n_heads: number of attention heads
    """

    self.attention = MultiHeadAttention(embed_dim, n_heads=8)
    self.norm = nn.LayerNorm(embed_dim)
    self.dropout = nn.Dropout(0.2)
    self.transformer_block = TransformerBlock(embed_dim, expansion_factor, n_heads)

def forward(self, key, query, x, mask):
    """
    Args:
        key: key vector
        query: query vector
        value: value vector
        mask: mask to be given for multi head attention
    Returns:
        out: output of transformer block
    """

    #we need to pass mask mask only to 1st attention
    attention = self.attention(x, x, x, mask=mask) #32x10x512
    value = self.dropout(self.norm(attention + x))

    out = self.transformer_block(key, query, value)

    return out

class TransformerDecoder(nn.Module):
def __init__(self, target_vocab_size, embed_dim, seq_len, num_layers=2, expansion_factor=4, n_heads=8):
    super(TransformerDecoder, self).__init__()
    """
    Args:
        target_vocab_size: vocabulary size of target
        embed_dim: dimension of embedding
        seq_len : length of input sequence
        num_layers: number of encoder layers
        expansion_factor: factor which determines number of linear layers in feed forward layer
        n_heads: number of heads in multihead attention
    """

    self.word_embedding = nn.Embedding(target_vocab_size, embed_dim)
    self.position_embedding = PositionalEmbedding(seq_len, embed_dim)

    self.layers = nn.ModuleList(
        [
            DecoderBlock(embed_dim, expansion_factor=4, n_heads=8)
            for _ in range(num_layers)
        ]
    )

    self.fc_out = nn.Linear(embed_dim, target_vocab_size)
    self.dropout = nn.Dropout(0.2)

def forward(self, x, enc_out, mask):
    """
    Args:
        x: input vector from target
        enc_out : output from encoder layer
        trg_mask: mask for decoder self attention
    Returns:
        out: output vector
    """

    x = self.word_embedding(x) #32x10x512
    x = self.position_embedding(x) #32x10x512
    x = self.dropout(x)

    for layer in self.layers:
        x = layer(enc_out, x, enc_out, mask)

    out = F.softmax(self.fc_out(x))

    return out

```

In [16]:

```
class Transformer(nn.Module):
    def __init__(self, embed_dim, src_vocab_size, target_vocab_size, seq_length, num_layers=2, expansion_factor=4, n_heads=8):
        super(Transformer, self).__init__()

        """
        Args:
            embed_dim: dimension of embedding
            src_vocab_size: vocabulary size of source
            target_vocab_size: vocabulary size of target
            seq_length: length of input sequence
            num_layers: number of encoder layers
            expansion_factor: factor which determines number of linear layers in feed forward layer
            n_heads: number of heads in multihead attention
        """

        self.target_vocab_size = target_vocab_size

        self.encoder = TransformerEncoder(seq_length, src_vocab_size, embed_dim, num_layers=num_layers, expansion_factor=expansion_factor, n_heads=n_heads)
        self.decoder = TransformerDecoder(target_vocab_size, embed_dim, seq_length, num_layers=num_layers, expansion_factor=expansion_factor, n_heads=n_heads)

    def make_trg_mask(self, trg):
        """
        Args:
            trg: target sequence
        Returns:
            trg_mask: target mask
        """
        batch_size, trg_len = trg.shape
        # returns the lower triangular part of matrix filled with ones
        trg_mask = torch.tril(torch.ones((trg_len, trg_len))).expand(
            batch_size, 1, trg_len, trg_len
        )
        return trg_mask

    def decode(self, src, trg):
        """
        for inference
        Args:
            src: input to encoder
            trg: input to decoder
        out:
            out_labels: returns final prediction of sequence
        """
        trg_mask = self.make_trg_mask(trg)
        enc_out = self.encoder(src)
        out_labels = []
        batch_size, seq_len = src.shape[0], src.shape[1]
        # outputs = torch.zeros(seq_len, batch_size, self.target_vocab_size)
        out = trg
        for i in range(seq_len): #10
            out = self.decoder(out, enc_out, trg_mask) #bs x seq_len x vocab_dim
            # taking the last token
            out = out[:, -1, :]

            out = out.argmax(-1)
            out_labels.append(out.item())
            out = torch.unsqueeze(out, axis=0)

        return out_labels

    def forward(self, src, trg):
        """
        Args:
            src: input to encoder
            trg: input to decoder
        out:
            out: final vector which returns probabilities of each target word
        """
        trg_mask = self.make_trg_mask(trg)
        enc_out = self.encoder(src)

        outputs = self.decoder(trg, enc_out, trg_mask)
        return outputs
```

```
In [17]: src_vocab_size = 11
target_vocab_size = 11
num_layers = 6
seq_length = 12

# Let 0 be sos token and 1 be eos token
src = torch.tensor([[0, 2, 5, 6, 4, 3, 9, 5, 2, 9, 10, 1],
                    [0, 2, 8, 7, 3, 4, 5, 6, 7, 2, 10, 1]])
target = torch.tensor([[0, 1, 7, 4, 3, 5, 9, 2, 8, 10, 9, 1],
                       [0, 1, 5, 6, 2, 4, 7, 6, 2, 8, 10, 1]])

print(src.shape, target.shape)
model = Transformer(embed_dim=512, src_vocab_size=src_vocab_size,
                    target_vocab_size=target_vocab_size, seq_length=seq_length,
                    num_layers=num_layers, expansion_factor=4, n_heads=8)
model

torch.Size([2, 12]) torch.Size([2, 12])
```

```
Out[17]: Transformer(
  (encoder): TransformerEncoder(
    (embedding_layer): Embedding(
      (embed): Embedding(11, 512)
    )
    (positional_encoder): PositionalEmbedding()
    (layers): ModuleList(
      (0-5): 6 x TransformerBlock(
        (attention): MultiHeadAttention(
          (query_matrix): Linear(in_features=64, out_features=64, bias=False)
          (key_matrix): Linear(in_features=64, out_features=64, bias=False)
          (value_matrix): Linear(in_features=64, out_features=64, bias=False)
          (out): Linear(in_features=512, out_features=512, bias=True)
        )
        (norm1): LayerNorm((512,), eps=1e-05, elementwise_affine=True)
        (norm2): LayerNorm((512,), eps=1e-05, elementwise_affine=True)
        (feed_forward): Sequential(
          (0): Linear(in_features=512, out_features=2048, bias=True)
          (1): ReLU()
          (2): Linear(in_features=2048, out_features=512, bias=True)
        )
        (dropout1): Dropout(p=0.2, inplace=False)
        (dropout2): Dropout(p=0.2, inplace=False)
      )
    )
  )
  (decoder): TransformerDecoder(
    (word_embedding): Embedding(11, 512)
    (position_embedding): PositionalEmbedding()
    (layers): ModuleList(
      (0-5): 6 x DecoderBlock(
        (attention): MultiHeadAttention(
          (query_matrix): Linear(in_features=64, out_features=64, bias=False)
          (key_matrix): Linear(in_features=64, out_features=64, bias=False)
          (value_matrix): Linear(in_features=64, out_features=64, bias=False)
          (out): Linear(in_features=512, out_features=512, bias=True)
        )
        (norm): LayerNorm((512,), eps=1e-05, elementwise_affine=True)
        (dropout): Dropout(p=0.2, inplace=False)
        (transformer_block): TransformerBlock(
          (attention): MultiHeadAttention(
            (query_matrix): Linear(in_features=64, out_features=64, bias=False)
            (key_matrix): Linear(in_features=64, out_features=64, bias=False)
            (value_matrix): Linear(in_features=64, out_features=64, bias=False)
            (out): Linear(in_features=512, out_features=512, bias=True)
          )
          (norm1): LayerNorm((512,), eps=1e-05, elementwise_affine=True)
          (norm2): LayerNorm((512,), eps=1e-05, elementwise_affine=True)
          (feed_forward): Sequential(
            (0): Linear(in_features=512, out_features=2048, bias=True)
            (1): ReLU()
            (2): Linear(in_features=2048, out_features=512, bias=True)
          )
          (dropout1): Dropout(p=0.2, inplace=False)
          (dropout2): Dropout(p=0.2, inplace=False)
        )
      )
    )
    (fc_out): Linear(in_features=512, out_features=11, bias=True)
    (dropout): Dropout(p=0.2, inplace=False)
  )
)
```

```
In [18]: out = model(src, target)
out.shape
```

```
Out[18]: torch.Size([2, 12, 11])
```

```
In [19]: # inference
model = Transformer(embed_dim=512, src_vocab_size=src_vocab_size,
                    target_vocab_size=target_vocab_size, seq_length=seq_length,
                    num_layers=num_layers, expansion_factor=4, n_heads=8)

src = torch.tensor([[0, 2, 5, 6, 4, 3, 9, 5, 2, 9, 10, 1]])
trg = torch.tensor([[0]])
print(src.shape, trg.shape)
out = model.decode(src, trg)
out

torch.Size([1, 12]) torch.Size([1, 1])
```

```
Out[19]: [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]
```

